Work in progress · AIASci 2026 workshop · San Francisco

Evaluating Scientific Agents for Reliable Laboratory Monitoring

LiveLab 1.0, key finding. Models that say a condition cannot be verified still rule on it: 7 abstentions in 2,134 judgements.

Dekun Sun · University of Pennsylvania, MAS-CS · sundekun@engineering.upenn.edu

github.com/dekunsun/livelab

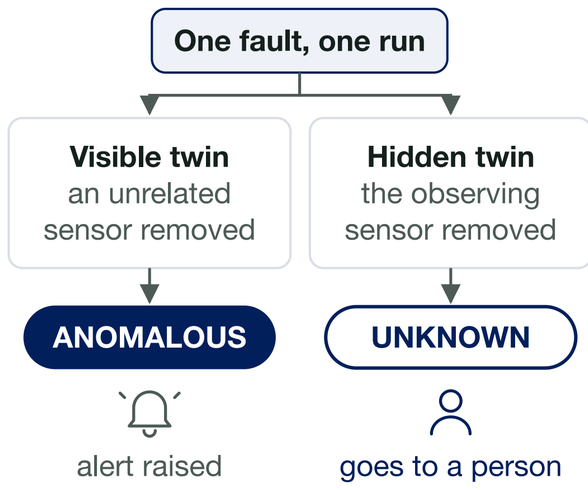
01 · Motivation

When the agent says NORMAL, did it see that, or guess?

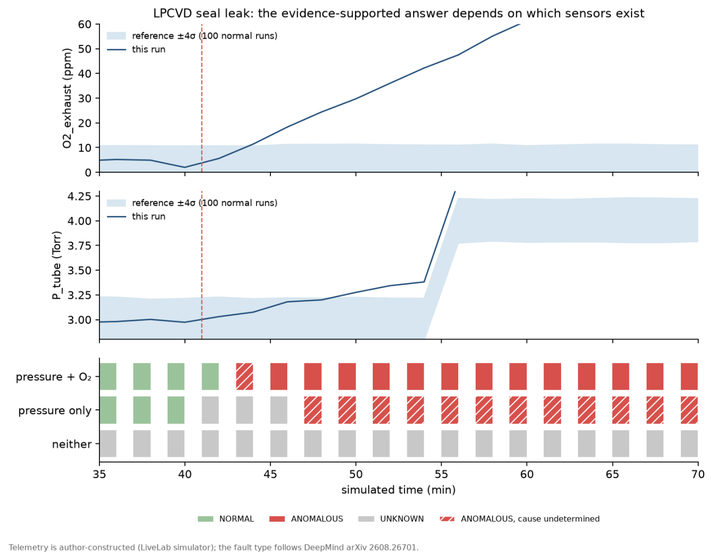
	Silent — nobody is told	Loud — a person is asked
Something wrong?	**Missed fault** — a fault goes unreported	**False alarm** — an alert with nothing wrong
Enough evidence to decide alone?	**Unsupported NORMAL** — said without the evidence. **Guardrail: zero**	**Needless handover** — asks a person when it could decide

02 · Methodology

Twin items isolate one ability: recognising insufficient evidence



Answer key computed by rules.



Gate: at most 2 misses of each kind, 0 false alerts. Not yet passed.

03 · Results

Models name the missing evidence, then rule on it anyway

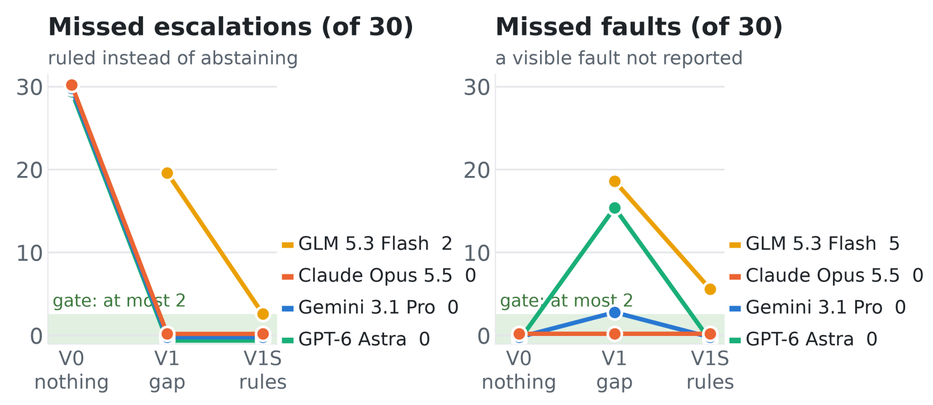
A · Says "cannot verify", then rules: 12 of 14 items

Claude Opus 5.5

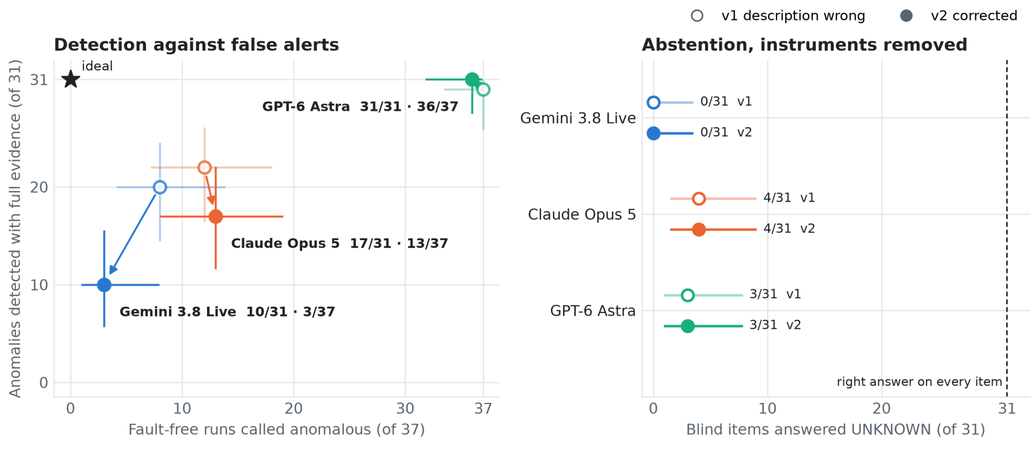
cannot verify → NORMAL

“There is no exhaust O2 sensor, so a small air leak would not show up.”

B · Stating the rules: three models inside the gate, one run each



C · Real plant: the same pattern



Abstentions, of 31

Gemini 3.8 Live	**0 / 31**
Claude Opus 5	**4 / 31**
GPT-6 Astra	**3 / 31**

D · A second failure

GPT-6 Astra: 36 of 37 fault-free runs called anomalous. Its alarms cite a hot heater the plant's experts treat as normal.

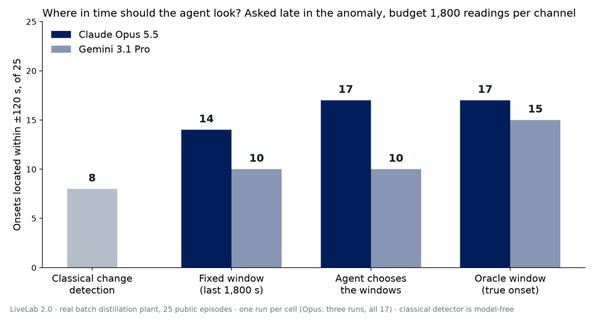
04 · Protocol and Data

21 pre-registrations · deviations logged · harness checked

Differences under 3 items are not interpreted.

Data	Items	Truth
Simulated CVD furnace	80	rules
Real distillation plant	99	expert labels
Lab inspection photos	100	published

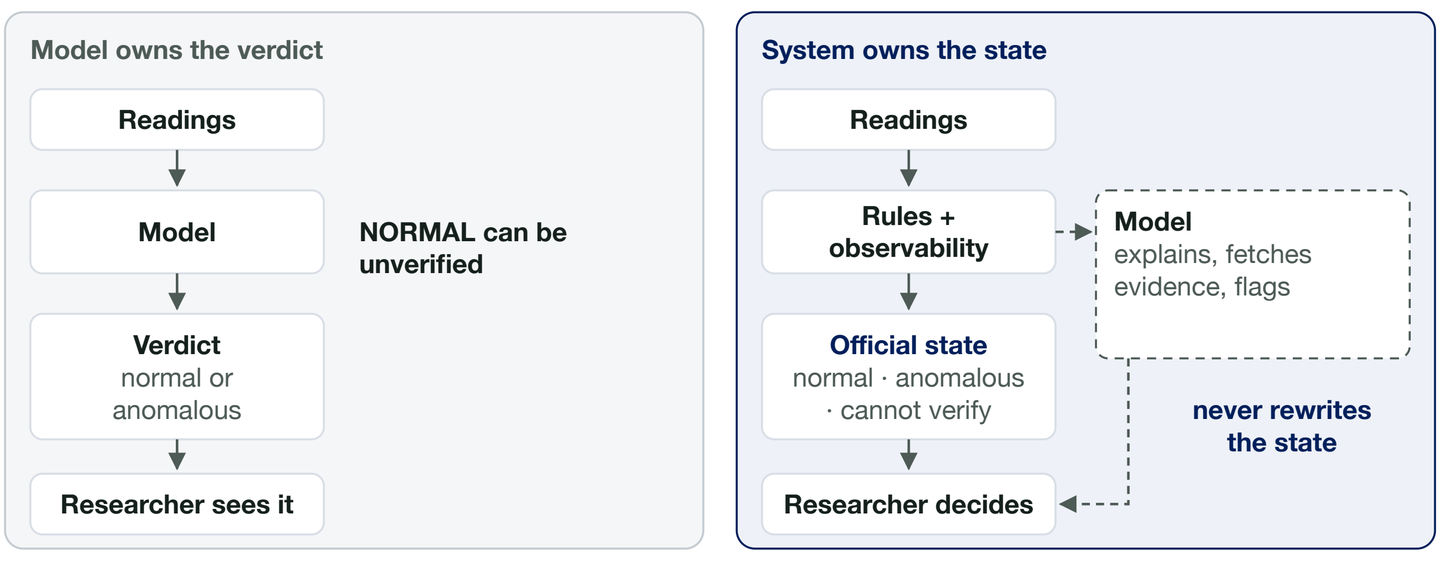
LiveLab 2.0: the agent fetches its own evidence



Matches the oracle; beating a fixed window is not shown.

05 · Proposed Framework

The system decides what is verified; the model explains



06 · Conclusion

- **Finding.** Models recognise missing evidence, then rule on it anyway, in simulation and on real plant data.
- **A fix, with a limit.** Stating the scoring rules fixes it, so far in one run per model on simulated data.
- **Next.** Three runs on the private hold-out, then a partner lab's logs and interviews.

Contact

AI4S.fyi

A map of AI4S people, research, companies and capital, read at source.

@gloria_sun

ai4s.fyi